

Towards Distributed Quantum Algorithms

Pablo Andres-Martinez



Master of Science by Research
Laboratory for Foundations of Computer Science
School of Informatics
University of Edinburgh
2018

Abstract

Acknowledgements

Declaration

I declare that this thesis was composed by myself, that the work contained herein is my own except where explicitly stated otherwise in the text, and that this work has not been submitted for any other degree or professional qualification except as specified.

(Pablo Andres-Martinez)

Table of Contents

1	Introduction	1
2	Quantum Computing: A brief overview	2
2.1	The principles of quantum computing	3
2.2	Building quantum computers	5
2.2.1	Scalability challenges	6
2.2.2	Models of computation	7
2.3	Programming on quantum computers	9
3	Distributed Quantum Computing	12
3.1	Communication through entanglement	12
3.1.1	Entanglement distillation	13
3.2	Distributing circuits	14
3.3	Distributed quantum architectures	15
4	Automated Distribution of Quantum Algorithms	17
4.1	Implementing non-local CNOT gates	17
4.2	Finding an efficient distribution	18
4.2.1	Vanilla algorithm	20
4.2.2	Bringing CNOT gates together	22
4.2.3	Adding the common-target method	23
4.3	Interchanging CNOT gates	28
4.4	An upper bound	29
5	Implementation details and Results	30
A	Hypergraph Partitioning	31
	Bibliography	32

Chapter 1

Introduction

Chapter 2

Quantum Computing: A brief overview

Quantum computing aims to take advantage of quantum mechanics to speed-up computations. There are many examples of problems that, although solvable by a standard (classical) computer, the time it would take to compute them is unreasonable in practice, regardless how large or fast your classical computer is. Some of these *intractable problems* can be solved efficiently on a quantum computer. This exhibits the vast gap in computational power between classical and quantum computers. Well known examples of such problems are:

- *Factorisation of large numbers*: In classical computers, all known factorisation algorithms take exponential time with respect to the input size. It is strongly believed that there is no way a classical computer can solve the problem efficiently – in fact, we are so confident about it that most widely used encryption systems, like RSA, rely on this. However, quantum computers are capable of solving the problem in polynomial time (i.e. making it tractable), using Shor’s algorithm (Shor, 1999).
- *Unstructured search*: The aim is to perform a brute-force search (i.e. requiring no prior knowledge about the search space) over N data-points. Classical computers have no other option than testing each data-point, so the time they take to perform the search is proportional to N . With a quantum computer, using Grover’s algorithm (Grover, 1996), the search is done in time proportional to \sqrt{N} .

Besides, in May of the present year, Raz and Tal (2018) gave formal proof of the existence of a large family of problems that a classical computer may never solve in polynomial time, but are solvable in polynomial time on a quantum one. Nevertheless,

all of these results are theoretical in nature, and giving experimental evidence of this gap in computational power is a highly active area of research, known as *quantum supremacy*.

Quantum computing would be very valuable in many areas of research that deal with problems that are intractable on classical computers. Some of the main applications that have been discussed in the literature are:

- *Chemistry, medicine and material sciences*: Calculating molecular properties on complex systems is extremely demanding for classical computers. However, polynomial algorithms for this problem are known for quantum computers (Lanyon et al., 2010).
- *Machine learning*: Finding patterns in a large pool of data is the essence of machine learning. Multiple quantum algorithms have been shown to be able to detect patterns that are believed not to be efficiently attainable classically (Biamonte et al., 2017).
- *Engineering*: Optimization and search problems are common in almost every area of engineering. Quantum computers are particularly well suited for these tasks, with Grover's algorithm being an obvious example.

For any of these applications we will require large scale quantum computers. Due to the obstacles in the way of building a large mainframe quantum computer (see §2.2.1), some authors have advocated the alternative of building a *quantum multicomputer*: a grid of small quantum computer units that cooperate in performing an overall computation (Van Meter et al., 2010). In the present work, particularly in Chapter 4, we contribute to this perspective, providing a method for efficiently distributing any quantum program originally designed for a monolithic quantum computer.

2.1 The principles of quantum computing

The advantages of using quantum mechanics to perform computations come down to the following three principles:

- *Superposition*: In classical computing, the unit of information is the *bit*, which may take one of two values: 0 or 1. In quantum computing, the *bit*'s counterpart

is the *qubit*, whose value may be *any linear combination* of the 0 state and the 1 state, known as a *superposition*, and usually written as:

$$|qubit\rangle = \alpha|0\rangle + \beta|1\rangle$$

where α and β are complex numbers that must satisfy: $\alpha^2 + \beta^2 = 1$.

A popular analogy of a qubit's superposition is a coin spinning¹: the classical states (0 and 1) are *heads* and *tails*, but when the coin is spinning, its state is neither of them. If we knew exactly how the coin was spinning, we would be able to describe the probability distribution of seeing heads or tails when it stops; these would be our α^2 and β^2 values. Besides, we may *measure* a qubit, and doing so corresponds in our analogy to abruptly stopping the coin, then checking if it is *heads* or *tails*. The essential aspect of this analogy is that, before measurement, the *qubit*'s state is neither 0 nor 1. Through certain operations – that would correspond to altering the axis of spin of the coin –, we may change the coefficients α and β of the superposition.

Interestingly, in quantum computing, we encode input and read output (after measurement) as standard classical binary strings, and thus, for input/output we use as many qubits as bits would be required. What superposition provides is the ability to – during mid-computation – maintain a superposition of all potential solutions to the problem, and update all of them simultaneously with a single operation to the qubits. In some sense, superposition allows us to explore multiple choices/paths of the computation, using only the resources required to explore a single one of those paths. And the number of paths we can explore simultaneously can be up to exponential, as a collection of N qubits may be in a state of superposition of all the possible 2^N classical states.

- *Interference*: As we just discussed, superposition gives us the ability to simultaneously explore different paths to solve a problem. However, in the end we will need to measure the qubits – stop the coins – and the result will be intrinsically random. For quantum computing to be any better than a probabilistic classical computer, we require the ability to prune the paths that have led to a result we do not want. This is precisely what *interference* provides: some operations on the qubits may make different classical states in the superposition cancel each other

¹ Note this is just an analogy, and while a coin spinning can be perfectly modelled using classical physics, a qubit can not.

out. Interference is at the core of any speed-up achieved by a quantum algorithm, and taking advantage of it is the main challenge when designing quantum algorithms.

- *Entanglement*: Quantum mechanics allows the possibility of having a pair of qubits a and b in a superposition such as:

$$|a, b\rangle = \frac{1}{\sqrt{2}}|0, 0\rangle + \frac{1}{\sqrt{2}}|1, 1\rangle$$

This implies that, when we measure the qubits, we may either read $a = 0, b = 0$ or $a = 1, b = 1$ as outcome, never $a \neq b$ (the coefficients for $|0, 1\rangle$ and $|1, 0\rangle$ are both 0). Then, what happens if we only measure a ? In this particular case, we would also know b 's outcome, without measuring it. In short, acting on one qubit has an instantaneous effect on the other. Whenever a group of qubits exhibits this property, we say they are *entangled*. Entanglement holds regardless how far apart a is from b ; for instance, they could be on two different quantum processing units of a distributed grid. Indeed, entanglement will be key in our discussion of distributed quantum algorithms, and we explain how to use it to perform non-local operations in §3.2.

2.2 Building quantum computers

Physicists have come up with different ways of realising qubits in labs. The key idea is to find a physical system that displays non-classical behaviour, and put it under the appropriate circumstances so we can manage its quantum properties, but noise in the environment may not interfere with these. Van Meter and Devitt (2016) give an excellent survey of the state of the art of quantum architectures. Among them, the three most developed are:

- *Quantum optics*: The state of a qubit is represented in the properties of photons, for instance, their polarization (Kok et al., 2007). A great advantage of this technology is that photons can be easily sent over long distances, while preserving the quantum state. Thus, protocols in quantum information that heavily rely on communication, such as Quantum Key Distribution (Shor and Preskill, 2000), are usually discussed and experimented using quantum optics. The downside of this technology is that it is very difficult to make photons interact, which is required for multi-qubit operations.

- *Ion-traps*: Each qubit is embodied as an ion, confined inside a chamber by means of an electric or magnetic fields. The qubit is acted upon by hitting the ion with electromagnetic pulses (e.g. laser light or microwave radiation). Groups of experimentalists have proposed how to scale up the technology (Weidt et al., 2016), and are currently building prototypes.
- *Superconductors*: Small circuits, similar to classical electrical circuits, are cooled down to near absolute zero so the quantum interactions of electrons are not obscured by other perturbations. Different parts of the circuit encode different qubits, which can be acted upon by applying electric potentials. One of the main advantages of this technology is that the technology required for building the chips is fairly similar to the one for classical computer chips. This technology seems to be the most experimentally developed. In fact, using it, both IBM and Intel have already built small generic-purpose quantum computers of 17-20 qubits.

However, for quantum computers to be useful in real world applications, their qubit count should raise up, at the very least, one order of magnitude. And, unfortunately, increasing the amount of qubits in a quantum computer is particularly difficult, due to some challenges we will now discuss.

2.2.1 Scalability challenges

There are two main challenges to overcome in order to build large scale quantum computers:

- *Decoherence*: In §2.1 we discussed the importance of having superposition in quantum computing, and we compared a qubit in superposition with a coin spinning. For similar reasons why a coin spinning will eventually stop, a qubit in superposition will eventually degenerate into a classical state (i.e. either $|0\rangle$ or $|1\rangle$): Physical systems have a tendency towards the state at which they are most stable, for the coin it is laying flat, for the qubit it is losing its superposition. This phenomenon is known as *decoherence*. Experimentalists attempt to increase the time it takes for the state of the qubit to degenerate². Decoherence constitutes

² A fundamentally different approach, *anyonic* (a.k.a. topological) quantum computing, has been proposed to avoid the problem of decoherence altogether, as it would use physical systems that, theoretically, can be completely protected against decoherence (Nayak et al., 2008). Although promising, currently this proposal has little experimental underpinning, and it is not regarded as attainable in the near future.

the main constraint to scalability of quantum computers, as it dictates the lifespan of qubits, limiting the number of operations that can be applied in a single program.

Certainly, the state of bits also degenerates in classical computers. However, in their case this is easier to account for: we can keep monitoring the bits, and make sure to correct any unwanted change. This is not so simple in quantum computers, as monitoring a qubit would require *measuring* it, and that destroys any quantum superposition. Nevertheless, it is still to some extent possible to protect our quantum state from errors – either due to decoherence or imperfect hardware – through quantum error correction routines (Knill et al., 2000). This is a very active area of research, and it will be key for the implementation of reliable large scale quantum computers.

- *Connectivity*: In order to run complex computations on qubits, we will need to be able to apply multi-qubit operations on any subset of the computer's qubits. However, it is not realistic to expect that quantum computers will have fast connectivity between all qubits, for instance due to spatial separation of these in the hardware. In classical systems, this problem is solved by a memory hierarchy, with a ceaseless flow of data going up and down of it, from main memory to registers (where computation is performed) and back. The memory hierarchy model works because data can stay idly in main memory while computation on the registers is carried out. However, in quantum computers we must avoid qubits being idle, as decoherence prevents the existence long-lasting memory. An alternative found in classical computers is to distribute the computation across different processing units each having its own local memory, which they use intensively, and communicating – through message passing – as little as possible. In §3.3, we discuss an abstract distributed quantum architecture in detail.

2.2.2 Models of computation

In this section, we give a brief introduction to some models of quantum computation relevant to the this thesis.

- *Circuit model*: Any operation on n qubits – as long as measurement (i.e. destruction of information) is not involved – can be represented as square matrix on complex numbers, of dimension 2^n . These matrices are always unitary, which

means that a matrix U satisfies $UU^\dagger = I = U^\dagger U$, where I is the identity matrix and A^\dagger is the conjugate transpose of A . Essentially, unitarity ensures that any operation on qubits can be reversed (i.e. undone), reason why this model is sometimes called the reversible model. Multiplying matrices AB corresponds to applying the operation described by B first, then A , on the same qubits. Application of two operations on disjoint set of qubits corresponds to the Kronecker product of the matrices $A \otimes B$. The fundamental concept is that any matrix can be represented as a product of other matrices, so we may decompose any operation into smaller building blocks: quantum gates.

Qubits are pictured as wires to which quantum gates are applied, similarly to a classical digital circuit. The set of quantum gates used is dependent on the architecture. There exist an (uncountable) infinite amount of different quantum operations, but a small finite set of them is enough to approximate any of them up to a desired error factor. The most common choice of such a universal gate-set is *Clifford+T*, which contains six one-qubit gates, and a single two-qubit gate. The depiction of the gates in that set, and some of their most important properties are shown in Figures **TODO**. Circuits are read from left to right.

TODO: A figure depicting, and giving the matrix of, each of CNOT, X, Y, Z, H, S and T. Figures showing $HXH = Z$; $XX = I$, $ZZ = I$, $YY = I$, $HH = I$; $SS = Z$; $TT = S$; H^2 CNOT $H^2 = \text{NOTC}$, and also their algebraic notation.

The CNOT gate (Figure ??) is particularly interesting. The qubit where the filled dot is acts as the ‘control’, and the qubit with \oplus acts as ‘target’. Whenever the control is $|0\rangle$, no change is made in either of the qubits; but if it is $|1\rangle$, an X gate (Figure ??) is applied to the target, flipping the state of the qubit. This works in any superposition, so if in

$$|c,t\rangle = \alpha|0,0\rangle + \beta|0,1\rangle + \gamma|1,0\rangle + \delta|1,1\rangle$$

$|c\rangle$ were acting as control and $|t\rangle$ as target, the outcome would be:

$$\text{CNOT} \cdot |c,t\rangle = \alpha|0,0\rangle + \beta|0,1\rangle + \gamma|1,1\rangle + \delta|1,0\rangle$$

TODO: A figure showing an abstract quantum circuit.

- *MBQC model:* The initials stand for Measurement Based Quantum Computing. Unlike the circuit model, where measurements are done at the very end of the

circuit, MBQC carries out computations by means of repeatedly measuring an initially entangled resource. The process, sketched in Figure ?? can be thought of as sculpting a rock. The initial entangled resource, which is a collection of entangled qubits forming a lattice structure, would be the rock. By measuring some qubits in the lattice – hitting the rock with a chisel – we remove some of the excess qubits, changing the overall state in the process. The outcome of measurements is probabilistic so, in order to provide deterministic computation, we must apply corrections on the neighbouring qubits whenever the measurement outcome deviated from the desired result. After multiple iterations of measurements and corrections, we end up with a set of qubits encoding the result. In this model, the input is incorporated into the lattice at the beginning of the process.

TODO: Figure from slides

In this way, any computation may be performed by applying 1-qubit measurements and 1-qubit correcting gates (controlled by classical signals). The initial resource state contains all the entanglement that is required, which may be prepared experimentally through multi-qubit interactions, such as Ising interactions (Raussendorf and Briegel, 2001), which are within our experimental capabilities. This model deals with the connectivity problem by applying a single operation involving *all* the qubits at the beginning of the process, then only requiring cheap single qubit operations for the rest of the computation. The main drawback of MBQC is the large amount of qubits that are required for even the simplest of operations. The MBQC model was proposed for the first time by Raussendorf and Briegel (Raussendorf and Briegel, 2001) under the name of *one-way quantum computer*, which highlights its main difference with the circuit (reversible) approach.

- *Distributed model:* We may find a balance between the circuit model and MBQC. In it, multiple small quantum processing units (QPUs) would run fragments of the overall circuit. Communication is achieved through a shared entangled resource, reminiscent of the MBQC approach. This model has been discussed in detail in the literature (Van Meter et al., 2010) and it is at the core of the main project from the Networked Quantum Information Technologies Hub (NQIT)³. In §3.3, we discuss an abstract distributed quantum architecture in detail.

³ A project supported by the UK National Quantum Technology program, aiming to provide scalable quantum computing.

2.3 Programming on quantum computers

As of today, most quantum programming languages are merely high level circuit descriptors: They provide the means to define circuits gate by gate, or build them up from combinations of smaller circuits. In this category fall all the well-known languages, such as *QCL* (Ömer, 2003) (imperative paradigm, and one of the first quantum programming languages ever implemented), *Q#* (Svore et al., 2018) (imperative, designed by Microsoft), and *Quipper* (Green et al., 2013) (functional, built on top of Haskell).

Besides, there are attempts at designing quantum programming languages that are completely hardware agnostic, meaning they aim to describe the computation, rather than a particular circuit that implements it. Examples of these are the different attempts at defining a quantum lambda calculus, for instance the ones by Van Tonder (2004) or Díaz-Caro (2017). However, these are not particularly programmer friendly, as they are generally quite verbose.

Most of the literature on quantum algorithms describes these by explicitly giving circuits that implement them. Fortunately, there is a constructive procedure, given by the Solovay-Kitaev theorem – of which Dawson and Nielsen (2005) give a good introductory review –, that takes any circuit and a choice of universal gate-set and outputs an efficient equivalent circuit using only those gates. Hence, programmers do not need to worry about the gates they are using when describing their circuits.

Unfortunately, the fact that algorithms are almost exclusively defined in the circuit model implies that other models of quantum computing are disregarded by a large portion of the community. In order to make other models of computation accessible, we need to provide automated procedures for transforming algorithms from the circuit model to these (and vice versa). Work has been done on the transformation from circuit to MBQC and backwards, the latter being the most challenging (Duncan and Perdrix, 2010). However, there is little amount of literature describing how to go from the circuit model to the distributed model. In §3.2 we give an overview of the existent work on that aspect, and identify the gap on the literature we aim to answer in this thesis.

Summary: Here are the key concepts to keep in mind while reading the rest of this thesis:

- Quantum computers provide a computing power well beyond the capabilities of classical computers, which would be exploitable in many areas of

science.

- Small quantum computers are already available.
- Scaling up is a challenging problem due to: *decoherence*, which may be overcome by the joint effort of error-correction, physics and engineering communities; and *connectivity*, which may be solved using distributed architectures.
- There is practically no programming support for distributed architectures.

Chapter 3

Distributed Quantum Computing

As we discussed in §2.2, there are different approaches on how to build quantum computers. Now that many of these have been experimentally demonstrated, having built small quantum computers, the question of how to scale up is increasingly relevant. This has led to the proposal of distributed architectures (Van Meter and Devitt, 2016).

In classical computing, an standard example of a distributed computer is the Non-Uniform Memory Access (NUMA) architecture: A system of independent computing nodes, each having its own local memory. In order to collaborate to perform an overall computation, the different nodes will need to communicate. In NUMA, they do so by accessing each other's memory. While nodes can manage their own local memory efficiently, accessing another node's memory is slow. Hence, we always attempt to minimise the amount of communication between nodes. A distributed quantum computer would follow the same principles, where each quantum processing unit (QPU) would own a collection of qubits (its local memory) and may access another QPU's qubits at the cost of some overhead, using entanglement.

3.1 Communication through entanglement

For a QPU to be able to access another's qubit, we must provide them with some sort of communication channel. Simply using a classical channel (sending bits) is not helpful: the whole point of representing the state of the computation in qubits is that they may be in a superposition of classical states, which would take up to an exponential amount of space and processing on classical bits. We could consider physically transporting the system that encodes the qubit from one QPU to another, and while that is certainly possible with photons, in general it is not feasible to have a channel that is both fast

and protects well against information loss due to decoherence.

In §2.1, we explained it was possible to affect a distant qubit by acting on another qubit with which it was entangled. We wish to exploit this property in order to allow a QPU to query another's QPU qubit. There are different levels of how strong a pair of qubits is entangled – intuitively, how much they affect each other. This is often formalised as the correlation between the qubits possible measurement outcomes; for instance, the pair of qubits $\frac{1}{\sqrt{2}}|0,0\rangle + \frac{1}{\sqrt{2}}|1,1\rangle$ is said to be *maximally entangled*, as the possible measurement outcomes are exclusively either $|0,0\rangle$ or $|1,1\rangle$, always matching in both qubits¹. Naturally, the most efficient communication channel will take advantage of entanglement in its strongest form, and so we will make use pairs of qubits entangled in this particular maximally entangled state. This qubit pair configuration is generally known as a Bell state, and Figure ?? shows how to prepare it.

TODO: fig:bell, show the circuit to make a bell state, how it is represented usually as a wire C and its Diract state.

An interesting property of the Bell state is shown in Figure ??: if a quantum gate, whose matrix representation is symmetric, is applied to one of the qubits, it is the same as if the gate was applied to the other qubit. In some sense, the gate can ‘slide’ through the entanglement, like beads on a string; as if the entangled state were a curved wire, connecting the pair of qubits. Hopefully, this serves as a first intuition of how Bell states are a natural choice for implementing quantum communication.

TODO: fig:sliding, show how a H gate can slide through an ebit.

So far, we have explained how to generate a Bell state inside a QPU (as in Figure ??). However, what we aim for is that two different QPUs each own one of the qubits from the Bell state. The challenge is then to send one of the qubits to another QPU, while preserving the entangled state. The problem of sharing a Bell state between two parties is solved by the *entanglement distillation* protocol (Bennett et al., 1996), which ensures the Bell pair is transmitted to their destination with an arbitrarily small error factor. We will now give a brief description of this protocol.

3.1.1 Entanglement distillation

TODO: This section

Sending an arbitrary quantum state across a channel is difficult, we could do it

¹ In total, there are four *maximally entangled* states of a pair of qubits: $\frac{1}{\sqrt{2}}|0,0\rangle + \frac{1}{\sqrt{2}}|1,1\rangle$ and $\frac{1}{\sqrt{2}}|0,0\rangle - \frac{1}{\sqrt{2}}|1,1\rangle$ give perfect correlation, while $\frac{1}{\sqrt{2}}|0,1\rangle + \frac{1}{\sqrt{2}}|1,0\rangle$ and $\frac{1}{\sqrt{2}}|0,1\rangle - \frac{1}{\sqrt{2}}|1,0\rangle$ give perfect anti-correlation

either fast or reliably, but not both. But if you do know what you are sending, you just send a lot of them.

Non-fidelity. If the ebit is not exactly in a state of the form $\alpha|0,0\rangle + \beta|1,1\rangle$ (i.e. it has a non-zero coefficient for $|0,1\rangle$ or $|1,0\rangle$)

Often in distributed quantum computing literature, a Bell state shared by two QPUs is known as an *ebit* (entangled-bit). We will also use this convention, and refer to each of the two qubits in the Bell state as the two *halves* of the ebit.

3.2 Distributing circuits

We now explain how ebits are used to allow a QPU to peek into another's qubits. Here, we will introduce the proposal of Yimsiriwattana and Lomonaco Jr (2004). Later on, in §4.1, we will extend this work with our own contributions.

We aim to split a given circuit and distribute the fragments across multiple QPUs. The gates that should operate over qubits on different QPUs are known as *non-local* gates. As we previously mentioned, any circuit can be converted to Clifford+T circuit (thanks to the Solovay-Kitaev theorem). In Clifford+T, the only gate that operates on more than one qubit is the CNOT. Hence, we only need to understand how CNOTs can be implemented non-locally.

TODO: A figure with a simple circuit and dashed lines identifying how we want to split it, and highlighting a non-local CNOT.

The construction we will use is a slight variation of what was proposed by Yimsiriwattana and Lomonaco Jr (2004), and it is shown in Figures **TODO**. In principle, we will use an *ebit* per non-local CNOT. We will call the QPU that holds the target qubit (the one with a \oplus) the ‘target QPU’ and similarly for the control qubit.

TODO: showing the overall scheme, with cat-entangler and cat-disentangler as black boxes.

The implementation of a local CNOT gate has three steps. First, we must apply what the authors refer to as the *cat-entangler* (Figure ??), which creates a local ‘copy’² of the control qubit inside the target QPU. In the process, the ebit half in the control QPU is measured (and thus destroyed), and the outcome is used to correct the other half, in the same spirit as in the MBQC model. Notice that the only informa-

² Note that there is no such thing as ‘copying’ a quantum state (due to the non-cloning theorem). What we mean here by ‘copying’ is generating, from $|\psi\rangle = \alpha|0\rangle + \beta|1\rangle$, the state $|\psi\rangle = \alpha|0,0\rangle + \beta|1,1\rangle$ which is fundamentally different from an actual copy: $|\psi, \psi\rangle = \alpha^2|0,0\rangle + \alpha\beta|0,1\rangle + \alpha\beta|1,0\rangle + \beta^2|1,1\rangle$.

tion physically crossing the boundary between blocks is the *classical* outcome of the measurement (a bit, either 0 or 1).

TODO: Show the circuit for the cat-entangler and the resulting state

Then, the CNOT gate is *applied locally* inside the target QPU, between its ebit half and the target qubit. At the end, the *cat-disentangler* must be applied (Figure ??), which simply destroys – with a measurement – the remaining ebit half and then corrects the control qubit, so the randomness of the measurement is counteracted. Once again, only classical information crosses the boundary.

TODO: Show the circuit for the cat-disentangler and the resulting state

In this way, we have implemented a non-local CNOT gate using one ebit and two classical bit messages between QPUs. However, the true advantage of this approach is attained when multiple non-local CNOTs are implementing using a single ebit: Once the cat-entangler is applied, any number of CNOTs whose target is in the same QPU, and that are controlled by the same qubit, may all be implemented by using the same ebit half as control, as shown in Figure ??.

TODO: Show the circuit for many CNOTs, now with cat-entangler and everything

Now, depending of how we choose to partition the circuit, there will be different groups of CNOTs that we may be able to implement using a single ebit. We will then wish to find the partition that requires the fewest ebits to implement all of its CNOT gates. This optimization problem is not discussed in the original paper, nor in any other work, as far as we know. It will be our main contribution in this thesis, along with an extension of the results just explained, both found in Chapter 4.

3.3 Distributed quantum architectures

A distributed quantum computer will have multiple quantum processing units (QPUs), each managing a small collection of qubits as their local memory. It should also have a subsystem specialised in generating and sharing ebits. As we discussed in §3.1.1, there is a compromise between the quality of the ebit – the quality of the communication channel – and the time it takes to prepare it. Fortunately, Cirac et al. (1999) showed that efficient distributed quantum computation using noisy ebits is feasible. Still, the generation of *ebits* remains the main bottleneck of the architecture, and thus we will want to minimise the required ebit count as much as possible.

TODO Detail the abstract distributed quantum architecture.

TODO Authors such as Van Meter et al. (2010) have discussed the experimen-

tal construction of a similar architecture. It is a multicomputer that uses photons for communication across QPUs, entangling stuff.

Chapter 4

Automated Distribution of Quantum Algorithms

4.1 Implementing non-local CNOT gates

In §3.2, we explained the proposal by Yimsiriwattana and Lomonaco Jr (2004) of how to implement a non-local gate. We will now extend their results.

In the original work, in order to implement multiple non-local gates using a single ebit, there must not be any operation on the control qubit between the non-local gates. However, many of the 1-qubit gates from the Clifford+T set commute with the CNOT gate. This means that, if there are operations in between CNOTs, we may transform the circuit to an equivalent version where the CNOT gates are brought together. And although some of the gates do not commute with CNOT, in some cases they can still be interchanged with it if an additional 1-qubit gate is added. All of the relevant circuit transformations are shown in Figures **TODO**. These can be checked by calculating the corresponding matrices of both sides and verifying they match.

TODO: Figures of how to push gates through both control and target

The second improvement comes by realising that the method used to implement multiple CNOT gates controlled by the same wire (explained in §3.2) can also be applied if multiple CNOT have a common target qubit instead. We will refer to the former method as the *common-control* method and *common-target* for the latter. The derivation of the common-target method is shown in Figure **??**, which uses some of the properties listed in §2.2.2. The main difference is that the CNOT itself is now applied in the control QPU instead of the target, so the cat-entangler and cat-disentangler must change accordingly.

TODO: CNOTtargetProof

4.2 Finding an efficient distribution

In this section we explain how we search for a suitable distribution of the circuit. But first, we establish what we mean by a distributed circuit to be efficient. This is based on what was discussed in §3.3 to be the bottle-necks of a distributed quantum architecture. A distributed circuit is efficient when there is:

- *Minimal amount of quantum communication* between the QPUs, meaning it requires as little number of ebits as possible. In comparison, message passing of classical bits is considered negligible and it is not taken into account.
- *Load-balance across the QPUs*, up to a tolerance margin. Our notion of load-balance is that the different QPUs have a similar number of qubits assigned to them. A uniform depth of the local circuits (i.e. length of the circuits) would be desirable. However, the distributed circuit depth is inherited from the original circuit, as none of our distribution techniques change the depth in a significant way. Hence, we will not take circuit depth into account, and instead assume that already known methods for depth reduction, such as the one described by da Silva et al. (2013), have already been applied on the input circuit, and may be applied again to each QPU's local circuit. As circuit depth is not something we aim to optimise, we consider the cost of local gates negligible.

The problem at hand is similar to the (k, ϵ) -graph partitioning problem. In it, a graph partition in k subgraphs has to be found, minimising the number of *cut edges*: edges that have their incident vertices in different subgraphs. Additionally, the resulting partition must satisfy that the number of vertices in each subgraph is less than $(1 \pm \epsilon) \frac{N}{k}$, where N is the total number of vertices in the graph. In Table ?? we list the correspondences between these two problems.

But there is a caveat. If we use graph partitioning naively, we will not be exploiting the fact that multiple CNOT gates may be implemented using a single ebit. In what follows, we will explain how to make use of *hypergraph* partitioning, instead of simple graph partitioning, to account for this aspect. A more detailed review of the hypergraph partition problem is given in Appendix A, here we summarise the key concepts:

Table 4.1: Correspondence between the graph partitioning problem and the efficient distribution of quantum circuits.

<i>Graph partitioning</i>	<i>Efficient distribution</i>
Vertices	Circuit wires
Edges	CNOT gates
Partitioned graph	Distributed circuit
Subgraph	QPU
Min. cut edges	Min. non-local gates
Uniform subgraph size	Load-balance

- Hypergraphs extend graphs to accommodate edges that may have more than two incident vertices. More formally, a hypergraph is a pair (V, H) , where V is the set of vertices and $H \subseteq 2^V$ is the collection¹ of hyperedges. Each hyperedge is represented as the subset of vertices from V it connects. We will not consider any notion of directionality.
- Hypergraph partitioning follows the same premise as graph partitioning. The user provides a hypergraph and two parameters (k, ϵ) , which have the exact same meaning as before. What the problem now attempts to minimise is a metric known as $\lambda - 1$, which is defined as follows: Given a partition of the hypergraph, the function $\lambda: H \rightarrow \mathbb{N}$ pairs each hyperedge with the number of different *blocks*² its vertices are in. Then, $\lambda - 1 = \sum_{h \in H} \lambda(h) - 1$ provides a measure of not only how many hyperedges are cut but also how many blocks are they connecting³.

In the following subsections, we explain how hypergraph partitioning can be used to find the best distribution of a circuit. First, we only use the implementation of non-local gates described by Yimsiriwattana and Lomonaco Jr (2004) reviewed in §3.2. Later on, we extend the algorithm to include the improvements we have proposed in §4.1.

¹ We will allow multiple hyperedges connecting the same vertices, in the same way as multigraphs allow multiple edges across any pair of edges.

² The term *block* is often used to refer to each of the sub-hypergraphs that comprise the hypergraph partition. It is the term we will use throughout this thesis.

³ Simply minimising the number of cut hyperedges is also an often used approach, but it is not as useful for our problem.

Algorithm 4.1: Builds the hypergraph of a given circuit. H may contain multiple hyperedges connecting the same vertices.

```

1  input: circuit
2  output: (V,H)
3  begin
4     $V \leftarrow \emptyset$ 
5     $H \leftarrow \emptyset$ 
6    hedge  $\leftarrow \emptyset$ 
7    foreach wire in circuit do
8       $V \leftarrow V + \{\text{wire}\}$ 
9       $H \leftarrow H + \{\text{hedge}\}$ 
10     hedge  $\leftarrow \{\text{wire}\}$ 
11     foreach gate in wire do
12       if gate == CNOT and controlOf(gate) == wire then
13         hedge  $\leftarrow \text{hedge} + \{\text{targetOf(gate)}\}$ 
14       else
15          $H \leftarrow H + \{\text{hedge}\}$ 
16         hedge  $\leftarrow \{\text{wire}\}$ 
17   end

```

4.2.1 Vanilla algorithm

The key challenge is how to use hyperedges to represent a collection of CNOT gates that, in case of being non-local, they could all be implemented using a single ebit. In this first version of the algorithm, we will group CNOTs together only if they have a common control wire, and there are no other gates in between their connections to that wire. We will create a *single hyperedge* for every such a collection of CNOT gates. The hyperedge's vertices correspond to the controlling wire and each of the different wires the CNOT gates target. Algorithm 4.1 builds the hypergraph in that way. It should be noticed that all CNOT gates from the input circuit are represented once and only once in the hypergraph.

We then solve the hypergraph partitioning problem (see Appendix A) on the resulting hypergraph. Once an efficient partition of the hypergraph is obtained, we map the partition back to the circuit, distributing it. The resulting partition is interpreted as follows in the circuit:

- The way the vertices are assigned in the blocks dictates how the corresponding wires are allocated to the different QPUs.

- Whenever a hyperedge has all of its vertices in the same QPU (i.e. it is not cut), all the CNOT gates it represents are implemented locally in that QPU, so no ebit is required (Figure ??a).
- If a hyperedge is cut once, some of the wires targeted by the CNOTs will appear in the same QPU as the control wire, so those CNOTs will be implemented locally. The rest of the target wires will be in a different QPU, so the corresponding CNOTs will have to be implemented non-locally, all using the same ebit (Figure ??b).
- In the case a hyperedge is cut more than once – its vertices are split among more than two blocks – we will proceed in the same manner, but now we will require multiple ebits (Figure ??c). As we previously defined, the function $\lambda: H \rightarrow \mathbb{N}$ tells us how many blocks a cut hyperedge connects. In order to connect $\lambda(h)$ blocks, we need at least $\lambda(h) - 1$ block-to-block connections: $\lambda(h) - 1$ ebits. The first halves of all of the ebits is kept by the QPU with the controlling wire, while the rest of the QPUs receive a single half each. This is why the metric we will need our hypergraph partitioner to minimise is $\lambda - 1$.

TODO Figure showing the different cases from the paragraph above, a, b and c. Give both the original circuit, the hypergraph and the black boxed distributed circuit (fig:vanillaCuts)

Following this interpretation of the hypergraph partition, we transform the original circuit to obtain its distributed version. To do so, for each cut hyperedge h , we insert $\lambda(h) - 1$ cat-entanglers just before the group of CNOT gates it represents; and $\lambda(h) - 1$ cat-disentangler right after them. Then, each of the CNOTs that have to be implemented non-locally are modified so their control wire is the corresponding QPU's local ebit half, while the target is unchanged.

TODO: A figure showing a simple circuit, its hypergraph and its distributed circuit (with cat-(dis)entangler as a box). It'd be great if the example on these three subsections was the same (fig:distribProcess)

A simple example of the distribution process is shown in Figure ??. The resulting circuit is a distributed version of the original one, that is efficient in the sense described in the beginning of this section. Each of the QPUs can be set to implement its own local circuit, distilling ebits and using them along classical communication whenever indicated.

4.2.2 Bringing CNOT gates together

In §4.1 we have shown that any 1-qubit gate in the Clifford+T set acting on the control wire of a CNOT gate, with the exception of the Hadamard gate, can commute with the CNOT up to some byproduct. Here we use this fact, applying some preprocessing on the input circuit that brings together nearby CNOT gates, allowing us to implement more non-local CNOT gates using a single ebit. Figure ?? gives an example of how these transformations – listed in Figures **TODO** – can lead to a more efficient distribution of the circuit.

TODO Figure with (hopefully) the same circuit as the previous figure, its version after CNOTpulling and its hypergraph partition, which should have less lambda-1 (fig:pulledCNOTexample)

The preprocessing procedure is fairly straight-forward: Exploring the circuit from left to right, whenever a CNOT gate is found, use the transformations listed in Figures **TODO** to move it as early in the circuit as possible. The procedure introduces some additional X gates, due to the transformations in Figures **TODO**. Fortunately, X is its own inverse (i.e. $XX = I$) and every 1-qubit gate in Clifford+T can be interchanged with X in a simple way (as shown in Figures **TODO**). Hence, we should not expect a significant increase in the depth of the circuit, as most byproduct gates will cancel each other out.

So far we have been talking about standard 1-qubit gates, but in practical circuits we are likely to find 1-qubit gates that are *classically-controlled*, meaning that a classical signal (a bit, either 0 or 1) decides whether the gate is applied or not. These are no issue for the distribution of the circuit, as this classical control may only require classical communication between QPUs. Concerning the preprocessing we just described, classically-controlled 1-qubit gates can commute with CNOT under the exact same circumstances as their uncontrolled version. The only difference is that, whenever a byproduct gate is created, we must make sure it is controlled by the same classical signal that controlled the original gate, as shown in Figure ??.

TODO Figure with an X gate classically controlled before the control of a CNOT. Then the resulting circuit after exchanging them (fig:classicalControl)

The same procedure can be used to commute 1-qubit gates across the target wire of the CNOT gates, although in this case, apart from H , S and T gates can not commute either. This additional preprocessing would have no effect at all on the vanilla version of the algorithm, but it will be beneficial after we apply our next extension, which

requires CNOT gates to be adjacent on the target wire.

4.2.3 Adding the common-target method

In §4.1 we showed that the trick for implementing multiple CNOT gates using a single ebit also works if they share a common target wire (instead of the control wire). This makes our optimisation problem more intricate: Before, for each CNOT gate we only had two options, either to make it local or non-local – whether to cut the hyperedges or not. But now, when the CNOTs are to be implemented non-locally, we can choose to implement them using the common-control or common-target method. In order to represent this choice in our hypergraphs, we will use two different kinds of hyperedges, whether what their CNOTs have in common is the control or the target. Figure ?? shows a simple circuit where neither of the options is a priori better.

TODO: Figure where with 3 CNOTs we show that depending on how they are grouped together, we may gain something or not, which ultimately depends on how the hypergraph is split. Include both hypergraphs (fig:BothEndsSimple) Caption: The hypergraph of the circuit built as in Algorithm 4.1 uses the common control of gates α and β , while the other one is built similarly, but on common target. Different line format for the hyperedges whether control or target.

Now, imagine the circuit from Figure ?? to be a fragment of a larger one. Then, it may happen that the rest of the circuit makes the most efficient distribution allocates wires A and B in one QPU and wire C in another. As shown in Figure ??, in this particular case using the common-target method saves us one ebit. Conversely, if A were the wire on a separate QPU and B, C were together, using the common-control method would be best. The conclusion is that the choice of control or target hyperedges on a particular fragment of the circuit is dependent on the overall partitioning.

Therefore, it would be best to build the hypergraph in such a way that the decision of using control or target hyperedges is not done a priori, but by the hypergraph partitioner itself. A naive approach, also shown in Figure ??, would be to include all of the hyperedges, both of control and target type, in the same hypergraph. But then either partitioning of the hypergraph cuts three hyperedges, so the hypergraph partitioner sees no difference between the two options, preventing it from taking into account this subtlety when optimising.

TODO: Figure of hypergraphs from prev Fig; with control hyps, target hyps, naively. The dashed lines represent two hypothetical ways of partitioning the hyper-

graph. Different line format for the hyperedges whether control or target (fig:BothEndsChallenge)

We propose to build the hypergraph as in Algorithm 4.2. This hypergraph requires an additional vertex per CNOT in the circuit, and the block where the hypergraph partitioner decides to include that vertex in dictates whether the CNOT is implemented as a control or target hyperedge. Figure ?? shows how the hypergraph is built step by step for our running example.

TODO: Figure where the hypergraph for the three CNOTs example is built step by step (as in the Algorithm above) (fig:BothEndsProcess)

The hypergraph built by Algorithm 4.2 has one caveat: When discussing load-balancing in §4.2, we explained that we were interested in allocating a uniform number of wires to each QPU. Previously, the hypergraph partitioner took care of this, as it tried to assign a uniform number of vertices to each block. But now, the hypergraph partitioner has no way of distinguishing between ‘wire’ vertices and ‘CNOT’ vertices, the latter being an artificial gadget that should not count towards load-balancing. The solution is simple, instead of the standard hypergraph partition problem, we apply a version of it where vertices can have a weight assigned (see Appendix A). Then, each wire-vertex is given weight 1, and each CNOT-vertex is given weight 0, effectively ignoring them for the load-balancing aspect.

We now explain how the hypergraph partition determines the distribution of the circuit:

- As in the vanilla version of the algorithm, the way wire-vertices are assigned to blocks determines how wires are allocated to QPUs.
- In the vanilla version, if a hyperedge was cut in $\lambda(h)$ blocks, we needed $\lambda(h) - 1$ ebits. The first halves of all of them were kept by the QPU with the control wire, while the rest of the QPUs received one ebit half each. Now, the exact same thing will happen for cuts on hyperedges of the control type, while in the case of target hyperedges, it is the QPU holding the target wire the one that keeps the one half of all the ebits.
- The block where a CNOT-vertex is assigned dictates in which QPU the CNOT operation itself is carried out. As we discussed in §4.1, if the non-local CNOT is implemented using the common-control method, the CNOT is applied in the target QPU, and vice versa. Hence, the block where the CNOT-vertex is assigned indicates the method that should be used to implement it.

Algorithm 4.2: Builds the hypergraph of a given circuit, without choosing whether CNOT gates are implemented through common control or common target.

```

1  input: circuit
2  output: (V,H)
3  begin
4    V  $\leftarrow \emptyset$ 
5    H  $\leftarrow \emptyset$ 
6    hedge  $\leftarrow \emptyset$ 
7    foreach wire in circuit do
8      V  $\leftarrow V + \{\text{wire}\}$ 
9      H  $\leftarrow H + \{\text{hedge}\}$ 
10     hedge  $\leftarrow \{\text{wire}\}$ 
11     hType  $\leftarrow \text{unknown}$ 
12     foreach gate in wire do
13       if gate == CNOT then
14         if controlOf(gate) == wire then
15           if hType == target then
16             H  $\leftarrow H + \{\text{hedge}\}$ 
17             hedge  $\leftarrow \{\text{wire}\}$ 
18             hType  $\leftarrow \text{control}$ 
19           if targetOf(gate) == wire then
20             if hType == control then
21               H  $\leftarrow H + \{\text{hedge}\}$ 
22               hedge  $\leftarrow \{\text{wire}\}$ 
23             hType  $\leftarrow \text{target}$ 
24             hedge  $\leftarrow \text{hedge} + \{\text{labelOf(gate)}\}$ 
25       else
26         H  $\leftarrow H + \{\text{hedge}\}$ 
27         hType  $\leftarrow \text{unknown}$ 
28         hedge  $\leftarrow \{\text{wire}\}$ 
29  end

```

An example is shown in Figure ???. Notice that the ebit required to implement each of the non-local CNOTs is guaranteed to exist as, by construction (Algorithm 4.2), each CNOT-vertex is always connected to exactly two wire-vertices, corresponding to its control and target wires. Whenever these three vertices are not assigned to the same QPUs, a hyperedge will be cut, and the required ebit will be generated.

TODO: Simple example of building step by step the previous hypergraph (fig:BothEndsDistrib)

Remark: It is easy to check that the distributed circuit we obtain using the vanilla algorithm (from §4.2.1) is the same as the one obtained by this approach, restricting it so either:

- a). The common-target trick is never used.
- b). CNOT gates are always executed in their target QPU.
- c). Our hypergraph partitioner never cuts target hyperedges.

Under this interpretation it is apparent that the problem of efficiently distributing a circuit using the common-control and common-target methods, and the problem of partitioning the hypergraph built by Algorithm 4.2, are indeed the same. All the relevant information from the circuit – all the possible choices of how to implement each an every CNOT in the circuit – is encoded in the hypergraph, while each aspect of the hypergraph partition can be mapped back to the circuit.

It is interesting to discuss the particular case in Figure ???. Here, a CNOT gate that must be applied non-locally between QPU 1 and QPU 2, ends up having its local fragment – the CNOT-vertex, and thus the CNOT gate itself – in neither of them, but in a distant QPU 3. This happens whenever QPU 1 and QPU 2 do not share the necessary ebit, but they are both sharing a suitable ebit with QPU 3. Intuitively, it can be seen as passing a message through a middle-man; sometimes, using communication through an already available middle-man is better than putting up a full blown connection just for a single message. The hypergraph partitioner acting on the hypergraph built by Algorithm 4.2 will automatically use this strategy whenever it reduces the number of ebits required. Notice, however, that this trick only allows to implement a single CNOT with both target and control wires in some particular QPUs. When the CNOT shares one of its wires with multiple CNOTs, it will often be best to use the common-control or common-target method to implement all of them together. This trick will be useful whenever there are stray CNOT gates that can not be efficiently grouped with other CNOTs.

TODO: The case where a non-local CNOT is actually applied in neither of the QPUs at its ends (fig:farCNOT)

In any case, some criticism must be considered here: if we put no constraint on the exploitation of ‘middle-men’ QPUs, it may happen that communication across many of the QPUs all use the same middle-man QPU to deliver their messages, potentially creating a bottle neck. We see two ways around this:

1. Accept that some circuits will naturally be prone to such a centralised communication network, where a few QPUs must be prepared to carry most of the communication work. This does happen in classical networking too, and sometimes the best option is a centralised system. As far as the hardware is designed with these needs in mind, centralisation may not be a bottle neck, but even an advantage.
2. In case really wish to have a decentralised network, you will need to find a way to ensure the hypergraph partition you get has some kind of load-balancing on communication. Fortunately, there is a simple way of load-balancing the usage of ebits: instead of giving CNOT vertices weight 0 as we previously discussed, we may give them some weight $\mu > 0$ that indicates how relevant communication load-balancing is in comparison to the uniformity of wire allocation across QPUs. Even better, we could use a custom version of the hypergraph partitioning problem where we provide three parameters instead of two, k, ϵ, η , where now ϵ acts as the tolerance for imbalance of wire-vertices, while η is a separate tolerance for imbalance of CNOT-vertices, both tolerances being enforced in any partition.

Across the figures of this subsection we have drawn differently hyperedges that group CNOTs with common-control or common-target. However, it should be noticed that the hypergraph partitioner does not use this information at any point. This distinction is made for explanation purposes, and it is only relevant when building the distributed version of the circuit. The procedure for distributing the circuit following a given hypergraph partition is very similar to that in the vanilla version of the algorithm (see §4.2.1). The only difference is that cat-entanglers and cat-disentanglers are slightly different whether the CNOT is implemented through common-control or common-target (see Figure ??).

As a wrap up of this section, in Figure ?? we show the same circuit being distributed in four different ways: with or without the extension from §4.2.2 and with

or without the extension from §4.2.3. This provides a simple example where both extensions are shown to reduce the number of ebits required to distribute a circuit.

TODO Figure example ‘interesting’ with different modes (fig:modes)

4.3 Interchanging CNOT gates

CNOT gates trivially commute when they are applied to different wires. When one of the wires is in common, if it has the same role (control or target) for both CNOTs, we can take advantage of it and implement them using a single ebit, as we just described. If both wires are in common and have the same role, the CNOTs cancel each other. But what happens if two CNOTs act on the same wire with different roles? In that case, we can still interchange the gates as in Figures **TODO**, but that creates additional CNOTs.

TODO: Interchange challenge (picture from phone) (fig:interchangeChallenge)

It may seem like preprocessing the circuit so it has the minimum possible number of CNOT gates would always be the best option for partitioning. However, this is not always true, as shown in Figure **??**. In some cases interchanging CNOTs may unlock a more efficient partitioning of the circuit, regardless of it adding more CNOTs. In other cases, it will have no benefit, and add an extra CNOT to take into account when partitioning. This compromise prevents us from knowing a priori the best choice. Providing the hypergraph partitioner with the flexibility of deciding either to interchange a particular pair of CNOTs or not – in the same spirit we did in §4.2.3 – would be the best solution. However, encoding such a choice in a hypergraph partitioning problem is very difficult, due to the following reasons:

1. *The way CNOT gates are ordered in the circuit is important:* This is something we omit in the hypergraph built by Algorithm 4.2; looking at the final hypergraph in Figure **??**, we can not tell whether α goes before or after β . This information is key when interchanging, as it will determine which are the new neighbours of the interchanged CNOTs. This could potentially be accounted for by imposing that any hyperedge is an ordered list of vertices⁴. However, the standard hypergraph partitioning problem does not take into account such ordering, so we would need to define a custom hypergraph partitioning problem where this information is somehow taken into account.

⁴ For instance, the first vertex corresponding to a circuit wire and the rest, corresponding to the different CNOT gates, ordered as in the circuit.

2. *Interchanging CNOT gates adds new CNOTs*: This problem is substantially different from the problem in §4.2.3, where the count of CNOTs remained the same, and thus the only degree of freedom was whether each CNOT was implemented as control or as target hyperedge. The essence of our solution in §4.2.3 is to represent all of the options in the hypergraph. However, if we were to interchange a pair of CNOTs, new choices would become available: Is it worth to interchange the CNOTs again with their new neighbours? Should the byproduct CNOT be itself interchanged further? Although the number of options to take into account is finite, it is considerably larger. And what is worse, each choice would not be independent from the rest – as some interchanges are only available if others have been done before – so the structure of the hypergraph would likely be quite complex to accommodate this aspect.

Instead of encoding this choice within the hypergraph partitioning problem, another option is to apply some postprocessing once the circuit has been partitioned. This postprocessing would exhaustively apply the transformations in Figures **TODO** to each pair of potentially interchangeable CNOTs, keeping those transformations that reduced the ebit count.

TODO Is this a greedy algorithm? (most probably, a greedy algorithm won't be optimal) Should I give more details on the strategy? Am I implementing this too?

4.4 An upper bound

TODO: Finally, for the sake of comparison, we will use some theoretical results on quantum circuit decomposition, in order to estimate an upper bound of the number of ebits needed to distribute any quantum process on N qubits.

TODO: Discuss structured vs unstructured as the reason why we expect our algorithm to perform better.

Chapter 5

Implementation details and Results

Appendix A

Hypergraph Partitioning

Bibliography

- Bennett, C. H., Brassard, G., Popescu, S., Schumacher, B., Smolin, J. A., and Wootters, W. K. (1996). Purification of noisy entanglement and faithful teleportation via noisy channels. *Physical review letters*, 76(5):722.
- Biamonte, J., Wittek, P., Pancotti, N., Rebentrost, P., Wiebe, N., and Lloyd, S. (2017). Quantum machine learning. *Nature*, 549(7671):195.
- Cirac, J., Ekert, A., Huelga, S., and Macchiavello, C. (1999). Distributed quantum computation over noisy channels. *Physical Review A*, 59(6):4249.
- da Silva, R. D., Pius, E., and Kashefi, E. (2013). Global quantum circuit optimization. *arXiv preprint arXiv:1301.0351*.
- Dawson, C. M. and Nielsen, M. A. (2005). The solovay-kitaev algorithm. *arXiv preprint quant-ph/0505030*.
- Díaz-Caro, A. (2017). A lambda calculus for density matrices with classical and probabilistic controls. In Chang, B.-Y. E., editor, *Programming Languages and Systems*, pages 448–467, Cham. Springer International Publishing.
- Duncan, R. and Perdrix, S. (2010). Rewriting measurement-based quantum computations with generalised flow. In *International Colloquium on Automata, Languages, and Programming*, pages 285–296. Springer.
- Green, A. S., Lumsdaine, P. L., Ross, N. J., Selinger, P., and Valiron, B. (2013). Quipper: a scalable quantum programming language. In *ACM SIGPLAN Notices*, volume 48, pages 333–342. ACM.
- Grover, L. K. (1996). A fast quantum mechanical algorithm for database search. In *ANNUAL ACM SYMPOSIUM ON THEORY OF COMPUTING*, pages 212–219. ACM.

- Knill, E., Laflamme, R., and Viola, L. (2000). Theory of quantum error correction for general noise. *Physical Review Letters*, 84(11):2525.
- Kok, P., Munro, W. J., Nemoto, K., Ralph, T. C., Dowling, J. P., and Milburn, G. J. (2007). Linear optical quantum computing with photonic qubits. *Reviews of Modern Physics*, 79(1):135.
- Lanyon, B. P., Whitfield, J. D., Gillett, G. G., Goggin, M. E., Almeida, M. P., Kassal, I., Biamonte, J. D., Mohseni, M., Powell, B. J., Barbieri, M., et al. (2010). Towards quantum chemistry on a quantum computer. *Nature chemistry*, 2(2):106.
- Nayak, C., Simon, S. H., Stern, A., Freedman, M., and Sarma, S. D. (2008). Non-abelian anyons and topological quantum computation. *Reviews of Modern Physics*, 80(3):1083.
- Ömer, B. (2003). *Structured quantum programming*. na.
- Raussendorf, R. and Briegel, H. J. (2001). A one-way quantum computer. *Physical Review Letters*, 86(22):5188.
- Raz, R. and Tal, A. (2018). Oracle separation of bqp and ph. In *Electronic Colloquium on Computational Complexity*.
- Shor, P. (1999). Polynomial-time algorithms for prime factorization and discrete logarithms on a quantum computer. *SIAM Review*, 41(2):303–332.
- Shor, P. W. and Preskill, J. (2000). Simple proof of security of the bb84 quantum key distribution protocol. *Physical review letters*, 85(2):441.
- Svore, K., Geller, A., Troyer, M., Azariah, J., Granade, C., Heim, B., Kliuchnikov, V., Mykhailova, M., Paz, A., and Roetteler, M. (2018). Q#: Enabling scalable quantum computing and development with a high-level dsl. In *Proceedings of the Real World Domain Specific Languages Workshop 2018*, page 7. ACM.
- Van Meter, R. and Devitt, S. J. (2016). Local and distributed quantum computation. *arXiv preprint arXiv:1605.06951*.
- Van Meter, R., Ladd, T. D., Fowler, A. G., and Yamamoto, Y. (2010). Distributed quantum computation architecture using semiconductor nanophotonics. *International Journal of Quantum Information*, 8(01n02):295–323.

- Van Tonder, A. (2004). A lambda calculus for quantum computation. *SIAM Journal on Computing*, 33(5):1109–1135.
- Weidt, S., Randall, J., Webster, S., Lake, K., Webb, A., Cohen, I., Navickas, T., Lekitsch, B., Retzker, A., and Hensinger, W. (2016). Trapped-ion quantum logic with global radiation fields. *Physical review letters*, 117(22):220501.
- Yimsiriwattana, A. and Lomonaco Jr, S. J. (2004). Generalized ghz states and distributed quantum computing. *arXiv preprint quant-ph/0402148*.